

ne-age-prediction-sequential-model

February 18, 2024

1 Abalone Age Prediction Sequential Model TensorFlow

Generally, the age of an Abalone is determined by the physical examination of the abalone but this is a tedious task which is why we will try to build a regressor that can predict the age of abalone using some features which are easy to determine. `## Importing Libraries`

```
[64]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

import tensorflow as tf
from tensorflow import keras
from keras import layers

import warnings
warnings.filterwarnings('ignore')
```

1.0.1 Loading Dataset

```
[65]: data = pd.read_csv("drive/MyDrive/abalone.csv")
data.head()
```

```
[65]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	\
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	

	Shell weight	Rings
0	0.150	15
1	0.070	7
2	0.210	9

3	0.155	10
4	0.055	7

```
[66]: # Shape
data.shape
```

```
[66]: (4177, 9)
```

```
[67]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Sex                    4177 non-null   object
1   Length                 4177 non-null   float64
2   Diameter               4177 non-null   float64
3   Height                 4177 non-null   float64
4   Whole weight           4177 non-null   float64
5   Shucked weight         4177 non-null   float64
6   Viscera weight         4177 non-null   float64
7   Shell weight           4177 non-null   float64
8   Rings                  4177 non-null   int64
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB
```

```
[68]: data.describe().T
```

```
[68]:
```

	count	mean	std	min	25%	50%	75%	\
Length	4177.0	0.523992	0.120093	0.0750	0.4500	0.5450	0.615	
Diameter	4177.0	0.407881	0.099240	0.0550	0.3500	0.4250	0.480	
Height	4177.0	0.139516	0.041827	0.0000	0.1150	0.1400	0.165	
Whole weight	4177.0	0.828742	0.490389	0.0020	0.4415	0.7995	1.153	
Shucked weight	4177.0	0.359367	0.221963	0.0010	0.1860	0.3360	0.502	
Viscera weight	4177.0	0.180594	0.109614	0.0005	0.0935	0.1710	0.253	
Shell weight	4177.0	0.238831	0.139203	0.0015	0.1300	0.2340	0.329	
Rings	4177.0	9.933684	3.224169	1.0000	8.0000	9.0000	11.000	

	max
Length	0.8150
Diameter	0.6500
Height	1.1300
Whole weight	2.8255
Shucked weight	1.4880
Viscera weight	0.7600
Shell weight	1.0050

Rings 29.0000


1.1 Exploratory Data Analysis

```
[69]: data.isnull().sum()
```

```
[69]: Sex          0
      Length      0
      Diameter    0
      Height      0
      Whole weight 0
      Shucked weight 0
      Viscera weight 0
      Shell weight 0
      Rings       0
      dtype: int64
```

Distribution of the data in male, female and infant

```
[70]: x = data['Sex'].value_counts()
      labels = x.index
      values = x.values
      plt.pie(values, labels = labels, autopct = '%1.1f%%')
      plt.show()
```



by the look of the above the pie chat shows that we have equal amount of data for male, female, and infant abalone

```
[71]: data.groupby('Sex').mean()
```

```
[71]:
```

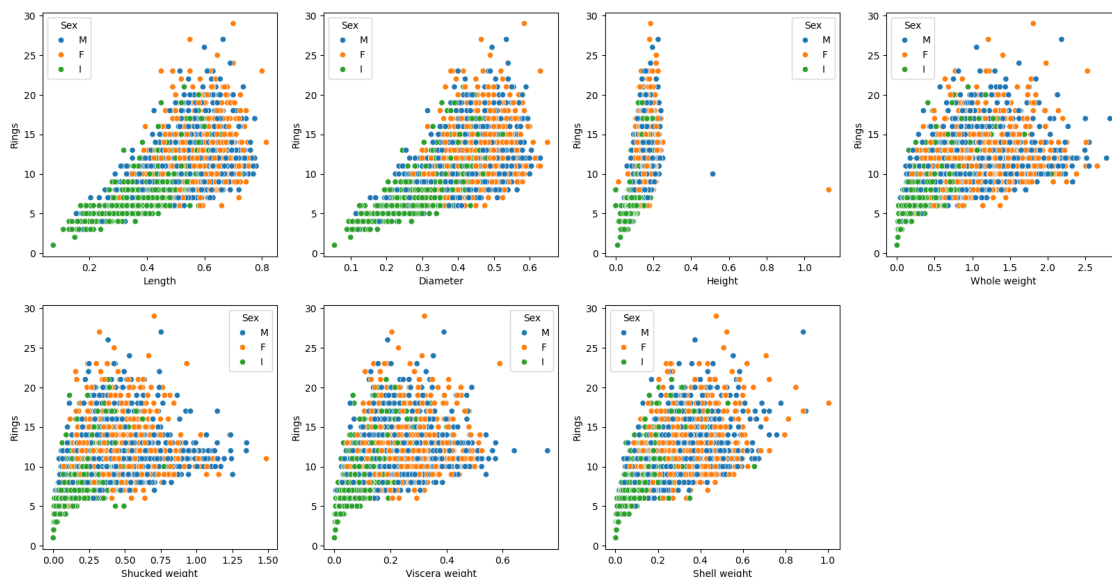
	Length	Diameter	Height	Whole weight	Shucked weight	\
Sex						
F	0.579093	0.454732	0.158011	1.046532	0.446188	
I	0.427746	0.326494	0.107996	0.431363	0.191035	
M	0.561391	0.439287	0.151381	0.991459	0.432946	

	Viscera weight	Shell weight	Rings
Sex			
F	0.230689	0.302010	11.129304
I	0.092010	0.128182	7.890462
M	0.215545	0.281969	10.705497

Here is an interesting observation that the life expectancy of the female abalone is higher than that of the male abalone. In the other features as well we can see that the height weight, as well as length in all the attributes of the numbers for female abalones, is on the higher sides.

```
[72]: features = data.loc[:, 'Length':'Shell weight'].columns
plt.subplots(figsize=(20,10))
for i, feat in enumerate(features):
    plt.subplot(2, 4, i+1)
    sns.scatterplot(data = data, x = feat, y='Rings', hue = 'Sex')

plt.show()
```



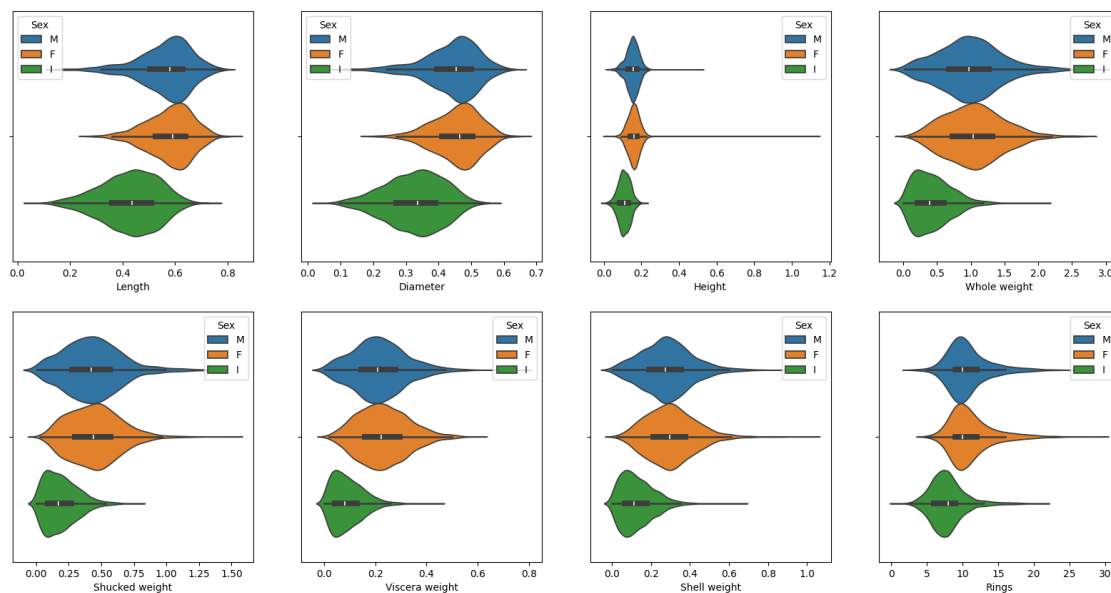
Observations from the above graph are as follows:

- A strong linear correlation between the age of the abalone and its height can be observed from the above graphs.
- Length and Diameter have the same kind of relation with age that is up to a certain age length increases and after that it became constant.

A similar kind of relationship is present between the weight and the age feature.

```
[73]: plt.subplots(figsize=(20, 10))
for i, feat in enumerate(features):
    plt.subplot(2, 4, i+1)
    sns.violinplot(data = data, x = feat , hue = 'Sex')

plt.subplot(2, 4, 8)
sns.violinplot(data = data, x = 'Rings', hue = 'Sex')
plt.show()
```



```
[74]: data.head()
```

```
[74]:   Sex  Length  Diameter  Height  Whole weight  Shucked weight  Viscera weight  \
0    M   0.455   0.365   0.095   0.5140      0.2245      0.1010
1    M   0.350   0.265   0.090   0.2255      0.0995      0.0485
2    F   0.530   0.420   0.135   0.6770      0.2565      0.1415
3    M   0.440   0.365   0.125   0.5160      0.2155      0.1140
4    I   0.330   0.255   0.080   0.2050      0.0895      0.0395
```

	Shell weight	Rings
0	0.150	15
1	0.070	7
2	0.210	9
3	0.155	10
4	0.055	7

Perform a One Hot Encoding on the Sex column.

```
[75]: data = pd.get_dummies(data, columns=['Sex'])
data.head()
```

```
[75]:   Length  Diameter  Height  Whole weight  Shucked weight  Viscera weight  \
0    0.455    0.365    0.095         0.5140         0.2245         0.1010
1    0.350    0.265    0.090         0.2255         0.0995         0.0485
2    0.530    0.420    0.135         0.6770         0.2565         0.1415
3    0.440    0.365    0.125         0.5160         0.2155         0.1140
4    0.330    0.255    0.080         0.2050         0.0895         0.0395
```

	Shell weight	Rings	Sex_F	Sex_I	Sex_M
0	0.150	15	0	0	1
1	0.070	7	0	0	1
2	0.210	9	1	0	0
3	0.155	10	0	0	1
4	0.055	7	0	1	0

Now I will separate the Features and target variables and split them into training and validation data.

```
[76]: features = data.drop('Rings', axis = 1)
target = data['Rings']
t='BY JosephWathome'
X_train,X_val, Y_train, Y_val = train_test_split(features, target, test_size = 0.2,
                                                random_state = 22)

X_train.shape, X_val.shape
```

```
[76]: ((3341, 10), (836, 10))
```

```
[77]: X_train.head()
```

```
[77]:   Length  Diameter  Height  Whole weight  Shucked weight  Viscera weight  \
3733    0.605    0.455    0.160         1.1215         0.5330         0.2730
3505    0.625    0.495    0.180         1.0815         0.4715         0.2540
3314    0.450    0.355    0.115         0.4385         0.1840         0.1080
1888    0.565    0.445    0.125         0.8305         0.3135         0.1785
3484    0.475    0.420    0.160         0.7095         0.3500         0.1505
```

	Shell weight	Sex_F	Sex_I	Sex_M
3733	0.2710	0	0	1
3505	0.3135	0	0	1
3314	0.1125	0	1	0
1888	0.2300	1	0	0
3484	0.1845	0	1	0

1.2 Model Architecture

I will impliment the **Sequential Model** that will contain the following parts: - I will have Two Connected Layers - I have included some **BatchNormalization** layers to enable stable and fast training and a **Dropout** layer before the final layer to avoid any possibility of overfitting.

```
[78]: model = keras.Sequential([layers.Dense(256, activation = 'relu', input_shape = [10]),
                                layers.BatchNormalization(),
                                layers.Dense(256, activation='relu'),
                                layers.Dropout(0.3),
                                layers.BatchNormalization(),
                                layers.Dense(1, activation='relu')])

model.compile(loss = 'mae',
              optimizer = 'adam',
              metrics = ['mape'])
```

```
[79]: model.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_12 (Dense)	(None, 256)	2816
batch_normalization_8 (Batch Normalization)	(None, 256)	1024
dense_13 (Dense)	(None, 256)	65792
dropout_4 (Dropout)	(None, 256)	0
batch_normalization_9 (Batch Normalization)	(None, 256)	1024
dense_14 (Dense)	(None, 1)	257

Total params: 70913 (277.00 KB)
Trainable params: 69889 (273.00 KB)
Non-trainable params: 1024 (4.00 KB)

converting the data to float32 for the model

```
[80]: X_train = X_train.astype('float32')
      Y_train = Y_train.astype('float32')
      X_val = X_val.astype('float32')
      Y_val = Y_val.astype('float32')
```

1.3 Model Training

```
[81]: history = model.fit(X_train, Y_train,
                          epochs=50,
                          verbose=1,
                          batch_size=64,
                          validation_data=(X_val, Y_val))
```

```
Epoch 1/50
53/53 [=====] - 2s 17ms/step - loss: 5.2891 - mape:
57.2654 - val_loss: 7.8946 - val_mape: 79.3217
Epoch 2/50
53/53 [=====] - 0s 7ms/step - loss: 4.1686 - mape:
47.0963 - val_loss: 7.1025 - val_mape: 71.5404
Epoch 3/50
53/53 [=====] - 0s 6ms/step - loss: 4.0295 - mape:
45.9125 - val_loss: 6.1870 - val_mape: 62.5931
Epoch 4/50
53/53 [=====] - 0s 7ms/step - loss: 3.9076 - mape:
44.6869 - val_loss: 5.1261 - val_mape: 51.8713
Epoch 5/50
53/53 [=====] - 0s 6ms/step - loss: 3.3706 - mape:
38.6901 - val_loss: 1.8799 - val_mape: 16.3903
Epoch 6/50
53/53 [=====] - 0s 7ms/step - loss: 2.7900 - mape:
31.8521 - val_loss: 1.8729 - val_mape: 19.3113
Epoch 7/50
53/53 [=====] - 0s 6ms/step - loss: 2.4032 - mape:
26.8200 - val_loss: 2.1171 - val_mape: 23.3601
Epoch 8/50
53/53 [=====] - 0s 7ms/step - loss: 2.2544 - mape:
25.0822 - val_loss: 2.0218 - val_mape: 22.1881
Epoch 9/50
53/53 [=====] - 0s 6ms/step - loss: 2.2369 - mape:
24.8090 - val_loss: 1.8889 - val_mape: 20.3759
Epoch 10/50
```


53/53 [=====] - 0s 6ms/step - loss: 2.1362 - mape: 23.4074 - val_loss: 1.8999 - val_mape: 20.6195
Epoch 11/50
53/53 [=====] - 0s 7ms/step - loss: 2.0768 - mape: 22.4461 - val_loss: 1.9617 - val_mape: 21.9786
Epoch 12/50
53/53 [=====] - 1s 10ms/step - loss: 1.9735 - mape: 21.5916 - val_loss: 1.6727 - val_mape: 17.3723
Epoch 13/50
53/53 [=====] - 1s 10ms/step - loss: 1.9031 - mape: 20.6304 - val_loss: 1.8582 - val_mape: 20.4082
Epoch 14/50
53/53 [=====] - 1s 10ms/step - loss: 1.8717 - mape: 20.0603 - val_loss: 1.6555 - val_mape: 17.2705
Epoch 15/50
53/53 [=====] - 1s 10ms/step - loss: 1.7624 - mape: 18.3546 - val_loss: 1.5901 - val_mape: 15.9443
Epoch 16/50
53/53 [=====] - 0s 9ms/step - loss: 1.7712 - mape: 18.3677 - val_loss: 1.6946 - val_mape: 17.7099
Epoch 17/50
53/53 [=====] - 1s 10ms/step - loss: 1.6906 - mape: 16.9950 - val_loss: 1.5600 - val_mape: 14.7417
Epoch 18/50
53/53 [=====] - 1s 10ms/step - loss: 1.6550 - mape: 16.5222 - val_loss: 1.6298 - val_mape: 15.6944
Epoch 19/50
53/53 [=====] - 1s 17ms/step - loss: 1.6752 - mape: 16.9274 - val_loss: 1.5083 - val_mape: 14.3708
Epoch 20/50
53/53 [=====] - 1s 16ms/step - loss: 1.6268 - mape: 16.3398 - val_loss: 1.7552 - val_mape: 16.7409
Epoch 21/50
53/53 [=====] - 1s 10ms/step - loss: 1.6223 - mape: 16.1914 - val_loss: 1.5222 - val_mape: 14.8033
Epoch 22/50
53/53 [=====] - 0s 8ms/step - loss: 1.5838 - mape: 15.9068 - val_loss: 1.5468 - val_mape: 15.1606
Epoch 23/50
53/53 [=====] - 0s 7ms/step - loss: 1.5931 - mape: 15.7406 - val_loss: 1.6474 - val_mape: 16.2875
Epoch 24/50
53/53 [=====] - 0s 7ms/step - loss: 1.6036 - mape: 15.8766 - val_loss: 1.6406 - val_mape: 16.1085
Epoch 25/50
53/53 [=====] - 0s 6ms/step - loss: 1.5829 - mape: 15.4957 - val_loss: 1.4986 - val_mape: 14.3542
Epoch 26/50

53/53 [=====] - 0s 7ms/step - loss: 1.5777 - mape: 15.5899 - val_loss: 1.5088 - val_mape: 14.2424
Epoch 27/50
53/53 [=====] - 0s 7ms/step - loss: 1.5696 - mape: 15.3485 - val_loss: 1.4703 - val_mape: 13.9026
Epoch 28/50
53/53 [=====] - 0s 7ms/step - loss: 1.5581 - mape: 15.3101 - val_loss: 1.4816 - val_mape: 13.9312
Epoch 29/50
53/53 [=====] - 0s 6ms/step - loss: 1.5493 - mape: 15.1414 - val_loss: 1.4758 - val_mape: 13.9861
Epoch 30/50
53/53 [=====] - 0s 7ms/step - loss: 1.5502 - mape: 15.0839 - val_loss: 1.5056 - val_mape: 14.3565
Epoch 31/50
53/53 [=====] - 0s 7ms/step - loss: 1.5452 - mape: 15.1466 - val_loss: 1.4956 - val_mape: 14.4843
Epoch 32/50
53/53 [=====] - 0s 7ms/step - loss: 1.5194 - mape: 14.9069 - val_loss: 1.5074 - val_mape: 13.8365
Epoch 33/50
53/53 [=====] - 0s 6ms/step - loss: 1.5433 - mape: 14.9819 - val_loss: 1.4817 - val_mape: 14.1558
Epoch 34/50
53/53 [=====] - 0s 7ms/step - loss: 1.5148 - mape: 14.9192 - val_loss: 1.4864 - val_mape: 14.3607
Epoch 35/50
53/53 [=====] - 0s 7ms/step - loss: 1.5150 - mape: 14.7361 - val_loss: 1.4988 - val_mape: 14.0730
Epoch 36/50
53/53 [=====] - 0s 6ms/step - loss: 1.5045 - mape: 14.6200 - val_loss: 1.5537 - val_mape: 15.3094
Epoch 37/50
53/53 [=====] - 0s 6ms/step - loss: 1.5367 - mape: 15.1302 - val_loss: 1.5128 - val_mape: 14.3585
Epoch 38/50
53/53 [=====] - 0s 7ms/step - loss: 1.5210 - mape: 14.8798 - val_loss: 1.4943 - val_mape: 14.1917
Epoch 39/50
53/53 [=====] - 0s 6ms/step - loss: 1.5033 - mape: 14.7799 - val_loss: 1.5056 - val_mape: 13.8984
Epoch 40/50
53/53 [=====] - 0s 7ms/step - loss: 1.5133 - mape: 14.8882 - val_loss: 1.4816 - val_mape: 13.8616
Epoch 41/50
53/53 [=====] - 0s 8ms/step - loss: 1.5069 - mape: 14.7457 - val_loss: 1.5079 - val_mape: 14.5612
Epoch 42/50

```

53/53 [=====] - 0s 8ms/step - loss: 1.5146 - mape:
14.8047 - val_loss: 1.5062 - val_mape: 14.5307
Epoch 43/50
53/53 [=====] - 1s 10ms/step - loss: 1.5073 - mape:
14.8130 - val_loss: 1.4719 - val_mape: 14.1408
Epoch 44/50
53/53 [=====] - 1s 10ms/step - loss: 1.5124 - mape:
14.8445 - val_loss: 1.4771 - val_mape: 13.9903
Epoch 45/50
53/53 [=====] - 1s 11ms/step - loss: 1.4892 - mape:
14.4589 - val_loss: 1.4877 - val_mape: 14.1112
Epoch 46/50
53/53 [=====] - 1s 10ms/step - loss: 1.4748 - mape:
14.3367 - val_loss: 1.4938 - val_mape: 14.4995
Epoch 47/50
53/53 [=====] - 1s 11ms/step - loss: 1.5032 - mape:
14.7512 - val_loss: 1.4663 - val_mape: 13.9024
Epoch 48/50
53/53 [=====] - 1s 10ms/step - loss: 1.4766 - mape:
14.3811 - val_loss: 1.5241 - val_mape: 13.8843
Epoch 49/50
53/53 [=====] - 0s 9ms/step - loss: 1.5035 - mape:
14.7135 - val_loss: 1.4785 - val_mape: 14.2238
Epoch 50/50
53/53 [=====] - 0s 6ms/step - loss: 1.4829 - mape:
14.3986 - val_loss: 1.4838 - val_mape: 14.3104

```

```

[82]: hist_df = pd.DataFrame(history.history)
      hist_df.head()

```

```


[82]:      loss      mape  val_loss  val_mape
0  5.289099  57.265411  7.894604  79.321724
1  4.168559  47.096336  7.102536  71.540443
2  4.029542  45.912514  6.187005  62.593052
3  3.907557  44.686943  5.126121  51.871315
4  3.370569  38.690086  1.879891  16.390261

```


```

[83]: hist_df['loss'].plot()
      hist_df['val_loss'].plot()
      plt.title('Loss v/s Validation Loss')
      plt.legend()
      plt.show()

```



```
[84]: hist_df['mape'].plot()  
hist_df['val_mape'].plot()  
plt.title('MAPE v/s Validation MAPE')  
plt.legend()  
plt.show()
```



From the above two graphs, we can certainly say that the two(mae and mape) error values have decreased simultaneously and continuously. Also, the saturation has been achieved after 15 epochs only.

By Joseph Wathome

```
[85]: print(t)
```

BY JosephWathome